

Lab 7: Linear Modeling

Introduction

In this lab we will learn how to work with linear models in RStudio.

Overall Goals

In this lab we will learn how to:

- add a data file to a project by importing from Excel,
- customize scatterplots,
- fit, visualize, and analyze linear models,
- and interpret residual plots.

Make Your Own Project in RStudio

As we did in the previous lab, start a new project:

- In RStudio, go to **File > New Project > New Directory > Empty Project**
- Make sure the parent directory (second textbox) is the folder where you want to keep your projects for this class. Use the **Browse** button if necessary.
- Enter a name for the new directory; it is good practice to avoid spaces in your file names by using underscores in place of spaces. For example, Lab_7 would be a good name.
- Click **Create Project**.
- In the Files pane, you should now see a file Lab_7.Rproj. This is the project configuration file, and you don't need to do anything with it.

Now you need to start a new RMarkdown report file.

- Go to **File > New File > R Markdown...**
- In the **Document** tab of the resulting dialog, give your report a title (Lab 7 Report) and put your name in the Author textbox. Keep HTML as the output format. Click **OK**.
- You should now see your new RMarkdown document at the upper left; **save** it. In the Save File dialog, enter the file name with no spaces: Lab7Report. (Do not use a filename extension.)
- Everything below the *first* provided code chunk is boilerplate and **should be removed**. Do so now.
- **Note:** The top-level section heading (one #) is already in use in this document for the report title. Use second-level headings (##) for the main sections of the report. (If you need subsections, use ###.)
- Use the **Insert** pulldown to add a new R code chunk. Add the command `library(hanoverbase)`. Use the chunk options dialog to disable warnings, disable messages, and “show nothing (run code)”.
- **Run the chunk** which you just created.

Import From Excel

In order to add a data file to the project itself, we start by uploading an Excel file:

- Download the Excel file that holds the relevant data: <https://hanoverstatslabs.github.io/resources/datasets/guns.xlsx>
- In the **Files** pane (lower right), click the **upload** button to start the Upload Files dialog.
- Navigate to your Downloads folder. Find the file `guns.xlsx` and click on it.
- Click **OK** to finish the upload. If this was successful, you should now see `guns.xlsx` in your Files pane.

Now we need to *import the actual data from the Excel file into the RStudio project* (both into the console and into the report):

- Click on the `guns.xlsx` file. Choose the option **Import Dataset**. This should bring up the **Import Excel Data** dialog.
- If the preview looks reasonable, click the **Import** button to do the import. You should now see the `guns` data in the data viewer, and a couple of lines of import code in the console. Next we will add these lines into our report.
- Make a new R code chunk in your report (below the `hanoverbase` chunk).
- There are two lines of code to copy from the console and paste into the new chunk. The first looks like `library(readxl)` and the second begins with `guns <-`. You do not need to run this chunk at this time, since we've already run those commands. But you need to have them in the chunk when you knit the report.

Explanation of the Variables in the Guns Dataset

We created the file `guns.xlsx` from data provided at www.openintro.org/stat/data. The five variables in the data are as follows:

country Name of the country.

mort_rate The number of gun-related deaths per 10,000 population.

own_rate The number of guns per 100 population. Note that this counts numbers of guns, not numbers of people. (Do you think it includes unregistered guns?)

hdi The country's numeric Human Development Index (a composite statistic of life expectancy, education, and per capita income indicators).

continent The continent that the country belongs to.

Statistical Investigations

You are now ready to start working on your report. The sections below give you questions to answer and commands to try. Here are a few reminders:

- **Important formatting note:** As you work through this lab assignment, answer the questions which are posed by typing into your R Markdown report, using formatting elements to make the report easy to read. See: Basic Cheatsheet¹
- **R Chunks:** As usual, create R Chunks (use the `Insert` pulldown) for your R commands. Use the R Cheatsheet for help as needed.
- **Knit Early and Often**

¹ [../cheatsheet.html](https://github.com/josiahelam/cheatsheet.html)

Warmup

To get warmed up and familiarize ourselves with the variables, we will start with some one-variable investigations.

1. Make a `favstats` summary and a histogram for the `own_rate` variable; you should adjust the number of breaks in order to get a good view.

Also make a labeled dotplot to show the countries in the dataset, sorted by gun ownership rate:

```
sum(~own_rate | country, data=guns) %>% sort() %>% dotplot()
```

If the dotplot looks too squeezed in the generated report, put it in its own chunk and use the chunk's options menu and the option **Use custom figure size** to specify the desired height for the chunk.

You should see two clusters in the dotplot, a couple of high outliers, and one extremely high outlier.

Describe the distribution of `own_rate`. What are the outliers and what are their gun ownership rates (number of guns per 100 population)?

2. Do the same for `mort_rate`. Describe what you find.
3. Do the same for `hdi`. Describe what you find.

Mortality and Gun Ownership

We wonder if countries with high gun ownership rates also have high gun-related mortality rates. We can investigate this with a scatterplot. We start by making a scatterplot for `mort_rate` (y) vs. `own_rate` (x). We attach the name `graph1` to the plot so that we can recall the plot in the future:

```
```r
graph1 <- xyplot(mort_rate~own_rate, data=guns)
graph1 # this is just to display the graph
```
```

4. Based on the scatterplot we just made, describe the overall pattern (including direction, if any) of the data and identify any unusual points.
5. We'll do some modeling without removing any observations to start with. Let's start with a smooth fit curve. Recall that to see all of the available color names in R, you can run the command `colors()` in the console.

```
ladd(panel.loess(x, y, col="your color here", lwd=2), plot=graph1)
```

Does it look like a linear model is a good fit for these data? Explain.

6. Let's see what a linear model looks like on this plot:

```
ladd(panel.lmline(x, y, col="your color here"))
```

In order to find the equation of our linear model, we need to calculate the slope and intercept. And to assess how well the model fits the data, we should calculate the square of the correlation (R-square linear).

```
fit1 <- lm(mort_rate~own_rate , data=guns)
coefficients(fit1)
r1 <- cor(mort_rate~own_rate , data=guns)
c("r"=r1 , "rsquared"=r1^2)
```

The last two numbers in the printout are the correlation r and its square r^2 respectively. The two previous numbers are the intercept and slope for the linear model.

- a. Write the equation of the linear model.
 - b. Assess how well the linear model fits the data. Explain.
7. In order to further assess the appropriateness of the linear model, we look at a residual plot, showing the **residual** (y minus fitted) vs. fitted. We add a horizontal line at 0 to help us judge the presence of a pattern:

```
xypplot(resid(fit1)~fitted(fit1))
ladd(panel.abline(h=0))
```

Remember that any remaining pattern in the residual plot indicates an incomplete model.

Do you see that there is a pattern in the residuals, or do they look “unpatterned” for the most part?

8. Correlation and regression are both susceptible to the effects of outliers and other influential points. Let’s see what happens when we filter out the U.S. entry from the data. The following command will create a new dataset called `gunsFiltered` by removing the row for the U.S. from the `guns` dataset. You will need to **replace the “..change this..”** with an expression involving the various variables, to leave out the U.S. (Hint: The U.S. is the only country with a very high gun ownership rate).

```
gunsFiltered <- guns %>% filter( ..change this.. )
```

If this was done correctly, you should see a new dataset entry in your environment that does not contain the U.S.

Repeat the code for #4 - #7 with the newly filtered dataset (`gunsFiltered`), changing the names `graph1`, `fit1` and `r1` to `graph2`, `fit2` and `r2` respectively. In the new scatterplot, for example, you should **no longer see** the point for the U.S. (`own_rate > 80`).

- a. Has your perception of the association changed? Explain.
 - b. Is there still a positive association? A linear association?
 - c. Has the correlation changed? If so, how?
9. Now we will show both linear models on the original (unfiltered) scatterplot. This demonstrates the effect of a single influential observation on the modeling process.

```
ladd(panel.abline(fit1 , col="black" , lwd=2), plot=graph1)
ladd(panel.abline(fit2 , col="magenta" , lwd=2))
```

Explain why the second line (`fit2`) has a smaller slope.

Other Relationships

10. We will now briefly take a look at how the two rates relate to the third variable of interest, hdi. We can do so by creating a **correlation table**, that shows the correlations between all possible pairs of selected variables. To make a correlation table, use the `cor` command in a new way:

```
cor(guns[,2:4]) # uses columns 2 - 4 in the data
```

Use a similar command to make a second table using the `gunsFiltered` data.

To see the correlation between two variables `x` and `y`, for example, find `x` in the columns and `y` in the rows. The number for that row and column gives the correlation for that pair.

- Explain why the correlations on the “main diagonal” will always be 1.0.
 - With which of the two rates is the hdi most strongly correlated?
 - hdi tries to measure how “advanced” the country is in terms of education, health, etc. How does it make sense that hdi is positively correlated with gun ownership rate?
11. We want to investigate whether the relationship between mortality rate and ownership rate varies depending on the hdi of the various countries. In order to do that, we will split the countries into groups based on their hdi. To find a good split, draw a histogram for hdi using a suitable number of breaks. Where would you split the values into different categories?
12. Create a new column `hdicat` in the dataset, which splits the cases depending on whether the hdi is lower or higher than 0.85. We pipe the hdi values into the `cut` command to create the desired breaks.

```
guns$hdicat <- guns$hdi %>% cut(breaks=c(0.75, 0.85, 0.95))
```

If this worked correctly, you should see a new column `hdicat` added to the `guns` dataset.

Now we make a paneled scatterplot for `mort_rate` versus `own_rate`, paneled by `hdicat`. We use `ladd` to add the linear models and `cor` to find the correlations:

```
xyplot(mort_rate~own_rate|hdicat, data=guns)
ladd(panel.lmline(x, y))
cor(mort_rate~own_rate, data=guns %>% filter(hdi <= 0.85))
cor(mort_rate~own_rate, data=guns %>% filter(hdi > 0.85))
```

What do you see in the resulting plot and correlations? Explain as best you can what is going on.

Submissions

- Make sure to knit one last time, then download the `Lab7Report.Rmd` file: in the **Files pane** (lower right), click the checkbox for the RMD file, then choose More > Export... > Download. *Do the same for Lab7Report.html.*
- Submit both files via Moodle.**