Advanced Lab 1: Data Preparation

Introduction

We can loosely split the "Data preparation" part of a project into three stages:

- Loading the data from a spreadsheet file or some other kind of storage. The readr and readxl packages help us with this stage.
- "Tidying" the loaded data to ensure the data is structured in a way that facilitates analysis. This mainly means that each variable is in one column. For example instead of having measurements over 5 days split across 5 columns, there should instead be one column of 5 times as many observations, with a second column identifying the day index for each. The tidyr package provides methods that help with this stage.
- Manipulating the data in various ways to obtain summaries or prepare for graphing. The dplyr package offers a rich set of methods to help with that.

All the packages mentioned above are part of the "tidyverse[1]" package collection, and are automatically included in our hanoverbase package.

Here is a summary of some of the key methods we will learn about in this lab:

read_excel	Read dataset from Excel spreadsheet
gather	Turn multiple columns to one, use a 2nd variable for the column names
separate	Split an entry into multiple parts and build different columns
rename	Change the name of a variable
mutate	Change or add a variable via computation using other variables
select	Restrict the dataset to a subset of the variables
filter	Only select certain rows based on some criterion
group_by	Group the rows of a dataset according to a factor variable
summarize	Produce a single value from each group in a grouped dataset
bind_rows	Bind datasets together by putting one below the other
recode_factor	Rearrange/rename a factor's levels
factor	Create a new factor or turn a character vector to factor

Importing

We start by importing the data from an external source. For this example, we will use the datasets/compression.xlsx[2] spreadsheet. It might help to view the spreadsheet in Excel.

In terms of providing us access to the data, the sheet is somewhat messy: It has the data spread over multiple tabs/sheets, and sheets often contain non-data components like summaries and graphs.

Before we start, make sure to create a new RStudio project and to start a new R Markdown document in that project, and add a library(hanoverbase) R-chunk in it.

---

[1] https://www.tidyverse.org/

[2] ../datasets/compression.xlsx

We will start by reading the data from the first sheet, titled "Upper Arm Circ". You will need to first have the dataset uploaded into your project directory. Use the "Upload" button in the "Files" pane for that. You can then click on the uploaded dataset to open up the import dialog. We will not fully use the import dialog, but it is a good starting point to look at the data.

In the import dialog, you can select the sheet you want, here the "Upper Arm Circ", from the "Sheet" drop down in the "Import Option" section. Then you should copy the couple of lines on the right, which provide syntax for the import. They will look something like this (make sure to simplify the path to only the last part, the actual file name):

```
library(readxl)
compression <- read_excel("compression.xlsx",
    sheet = "Upper Arm Circ")
```

Now, paste these lines in an R-chunk after canceling the import dialog. Remember that the first line loads a package for us, and only needs to be present once in the RMarkdown file.

Before running the chunk, we need to change the second command, to make sure we only try to load the first 8 data points, and also to improve the stored data set name:

```
compression.uac <- read_excel("compression.xlsx",
    sheet = "Upper Arm Circ",
    range = "A1:O9")
```

Notice the addition of the range="A1:O9" parameter, to control the range that will be imported. We learned about this extra parameter by reading the documentation for read_excel.

Let us also load another one of the sheets (you will need to fix the file path to match yours):

```
compression.lac <- read_excel("compression.xlsx",
    sheet = "Lower Arm Circ",
    range = "A1:O9")
```

This shows the basics of importing data from Excel. You can explore the readr and readxl packages to find out more. You start the exploration of an installed package by running help(package="readr"). This will give you a list of the functions in the package but also a link to user guides and "vignettes", which are basically tutorials contained within the package. You can see all vignettes by typing vignette(), and you can open a specific vignette by calling the vignette function with the vignette name, like so: vignette("readr")

Also check out the haven package for loading SPSS, Stata and SAS data directly. There are also ways to access a database directly; the package dbplyr[3] might be a good place to start for that.

Tidying

We will take various steps to "tidy" the data. Before we do that, let's discuss the data a bit. At this point you should use View in the console to view the two datasets we imported.

- The first column represents the subject ids, and each row of data corresponds to a subject.
- The "CS…" variables correspond to the subject using a "Compression Sleeve" while the "NT…" variables correspond to the subject receiving "Pneumatic Compression" therapy.

---

[3]https://dbplyr.tidyverse.org/articles/dbplyr.html

- The "…UAC…" part of the variable names represents that these are "Upper Arm Circumference" measurements.
- Finally, the remaining bits "PR, PO, 1, 2, 3, …" represent the time variable (pre-exercise, post-exercise, one day later, etc).

The data is in what is known as a "wide" form: In a "tidy" version of the data the kind of treatment, the quantity being measured, and the time of observation would all be three separate factor variables and we would have a single column for the measurement. In effect all the measurements currently present in the sheet will be placed in one column, and various information that is currently coded in the variable names will instead be stored in different variables/columns specifically created for that purpose. For instance there will be one variable to distinguish whether the measurement was done during the Compression Sleeve phase or during the Pneumatic Compression phase, another variable to describe what was measured (upper arm circ, lower arm circ etc), and another variable to distinguish which stage of a phase we are in (pre/post/day 1 etc).

We will now convert the data sheet into this form. Here as well as in many places later on we will be building a command iteratively, adding once step at a time. What you see in each of the following couple of code chunks is how the commands looks as it evolves, and therefore they will all correspond to the same chunk in your document, as it changes over time.

The first step in converting the data to a tidy form is to isolate the CSU columns as well as the subject column. We will use the select method for that, which picks out a subset of the variables. There are various ways to specify which variables to select, and you can look at the documentation or cheatsheets for more information.

```
compression.uac %>%
    select(X__1, CSUACPR:CSUAC5)
```

We could have stored this result in a dataset, but we will instead add another process step in the pipeline, which does the transformation. This step is called gather, and it simply collects some columns and puts them together into one column (the –X__1 tells it to combine all the other columns except the one named X__1). We also stored the result in a variable called sleeve.uac:

```
sleeve.uac <- compression.uac %>%
    select(X__1, CSUACPR:CSUAC5) %>%
    gather(key="time", value="value", -X__1)
View(sleeve.uac)    # Do in console
```

Notice how the new time column contains the variable names, and we now have 7 times as many rows, one for each combination of a subject and a variable.

Next we need to split the variable time into three parts:

- The fact that it is a "compression sleeve" treatment
- The fact that it is the upper-arm
- The timing (pre, post, day 1 etc)

The separate method helps us with this. Its syntax will look a bit weird. The second argument (into=) specifies the names to give to the three parts, while the third argument (sep=) specifies the cutoff points (after the first two letters, after the first five letters). This effectively splits each value like "CSUAPR" in three parts: the first two letters, the next three letters, and the rest.

```
sleeve.uac2 <- sleeve.uac %>%
        separate(time,
                  into=c("treatment", "measurement", "time"),
                  sep=c(2,5))
View(sleeve.uac2)
```

Finally, we will fix the first column. First we want to change its name to subject, with the method rename:

```
sleeve.uac2 %>%
    rename(subject=X__1)
```

And then change the character string values of the subject variable (like "sub 1") into the numbers themselves, using mutate. For this to work, we also need to load a string-manipulation library called stringr:

```
library(stringr)
sleeve.uac3 <- sleeve.uac2 %>%
    rename(subject=X__1) %>%
    mutate(subject = str_extract(subject, "\\d+"))
View(sleeve.uac3)    ## Do in the console
```

Creating a custom function for repeatable steps

We will likely need to repeat all these steps for each part of the dataset, and they will be more or less the same steps, so let's see if we can make a function out of it.

In order to make a function, we need to determine what it is that may be changing each time we do this. These are then "parameters" for this function. In our case the one thing that is of course changing every time is the data frame, and we will use the df parameter name to refer to it. But if we look at some of the variable names in the other sheets, we notice that the middle part of the name sometimes uses two letters and sometimes three. So this would be another parameter we can provide, called for example middleLetters. We will call our function convert.to.long:

```
convert.to.long <- function(df, middleLetters) {
    df %>%
        gather(key="time", value="value", -X__1) %>%
        separate(time,
                  into=c("treatment", "measurement", "time"),
                  sep=c(2, 2 + middleLetters)) %>%
        rename(subject=X__1) %>%
        mutate(subject = str_extract(subject, "\\d+"))
}
```

The \\d+ part is a regular expression[4] which tells R to match a sequence of one or more digits.

Then we could simply have done the following:

```
sleeve.uac <- compression.uac %>%
    select(X__1, CSUACPR:CSUAC5) %>%
    convert.to.long(3)
```

---

[4]https://stringr.tidyverse.org/articles/regular-expressions.html

Remember that we use the pipe %>% this fills in the first parameter, so we only had to provide the second parameter, middleLetters. Go ahead and replace the work you did with sleeve.uac, sleeve.uac2 and sleeve.uac3 with this single chunk above.

And if the variable names we wanted to use had a different number of "middle letters" we could provide that as a parameter. We can also set a "default value":

```
convert.to.long <- function(df, middleLetters=3) {
    ...
```

Then, if we do convert.to.long(), it will have the same effect as convert.to.long(3).

We will now do the same for the other part, the NT columns:

```
pneumatic.uac <- compression.uac %>%
    select(X__1, NTUACPR:NTUAC5) %>%
    convert.to.long(3)
```

Repeat this for the "lower-arm-circumference" variables:

```
sleeve.lac <- compression.lac %>%
    select(X__1, CSLACPR:CSLAC5) %>%
    convert.to.long(3)
pneumatic.lac <- compression.lac %>%
    select(X__1, NTLACPR:NTLAC5) %>%
    convert.to.long(3)
```

Binding data sets together

Now we have four data sets that all look similar in terms of the columns they contain etc. We would like to merge them together by simply placing them below each other and matching the columns. We can use the bind_rows method for this:

```
alldata <- bind_rows(sleeve.uac, pneumatic.uac, sleeve.lac, pneumatic.lac)
View(alldata)    # Do in console
```

There are many tools that help with more complicated data merges; look at the dplyr cheatsheet for some examples.

Handling factor variables

Now it is time to more properly code our factor variables, using mutate. In R, there are different *modes* for vectors/variables. One mode is *numeric* for scalar variables, another is *character* for strings. There is also *factor* which is what should be used for factor variables. A factor typically uses numbers for the underlying codes/levels, and it also contains a *value label* for each code/level. For example, take the treatment variable and turn it into a factor:

```
alldata$treatment
alldata$treatment %>% factor()
```

If you compare the outputs above, you will notice the results looking a tad different when we used the factor function. The important thing to know is that in order to do some of the standard statistical analysis steps, we need variables to be coded as "factors".

As we may have briefly mentioned brefore, the dollar sign is used to pick out a column from a dataset. As an alternative, you may also use the pluck command, part of the purrr package:

```
alldata %>% pluck("treatment") %>% factor()
```

Now let us code our categorical/factor variables as factors. There are two methods that can help with that. factor can be used if we are happy with the current value labels used, while recode_factor allows us to provide custom value labels. We use both of these methods in the code chunk below:

```
alldata <- alldata %>% mutate(
    subject = factor(subject),
    treatment = recode_factor(treatment, CS="Compression Sleeve", NT="Pneumatic Compression")
    measurement = recode_factor(measurement, UAC="Upper Arm Circ", LAC="Lower Arm Circ"),
    time = recode_factor(time, PR="Pre", PO="Post", `1`="Day 1", `2`="Day 2",
                               `3`="Day 3", `4`="Day 4", `5`="Day 5")
    )
```

The backticks around the numerical levels 1, 2 etc may seem weird. Whenever a value needs to be used on the left-hand-side of an equals sign, and it is not a "proper name" (meaning starts with a letter and has no spaces or other punctually), we need to enclose it in backticks. We will use this approach again later on when we have a situation with one of the names being the two words life expectancy.

We are now ready to do some work with this data. Here are two examples of graphs we can produce; you'll see more about this sort of graph in the graphing lab[5].

```
ggplot(alldata) +
    aes(x=time, y=value, group=subject, color=subject) +
    geom_path() +
    facet_grid(measurement~treatment)

ggplot(alldata) +
    aes(x=time, y=value, group=treatment, col=treatment) +
    stat_summary(geom="pointrange", fun.data=mean_se, fatten=3, alpha=0.6,
        position=position_dodge(width=0.1)) +
    geom_line(stat="summary", fun.y=mean, position=position_dodge(width=0.1)) +
    facet_grid(measurement~.)
```

Further work and transformations

Let us move on to doing some numerical work with this data. For instance we may want to compute the means and standard deviations for each day, for each treatment and for each type of measurement. In order to perform such computations, we need to first "group" the data by the grouping variables we want to use. The group_by method helps us with that. Then we can use the summarize function to produce numerical summaries per group.

```
summaries <- alldata %>%
    group_by(treatment, measurement, time) %>%
```

---

[5]LabIntroggplot.html

```
        summarize(mean=mean(value),
                  sd=sd(value),
                  median=median(value),
                  q1=quantile(value, 0.25))
summaries
```

We could then visualize those summaries in some way:

```
ggplot(summaries) +
    aes(x=time, y=mean, color=treatment, group=treatment) +
    geom_line() + facet_grid(measurement~.)
```

Let us try a more complicated transformation. Since different subjects start from different values at the pre-stage (you may have noticed that in the first graph you created earlier), we could try to normalize all the measurements for a given subject relative to the initial amount for that subject. So what we will do is the following:

- We will group our data by subject, treatment and measurement. Then each group contains the 7 measurements for a particular subject, during a particular treatment and a particular kind of measurements.
- Within each group, we can use the first function to get the first value, which should be the Pre-treatment value. We then want to divide all the values by it (we could instead have subtracted). We will use this transformation to create a new variable, called relChange:

```
relativedata <- alldata %>%
group_by(subject, treatment, measurement) %>%
    mutate(relChange=value/first(value))
```

We will also want to remove the "pre" entries, as they are no longer informative:

```
relativedata <- alldata %>%
group_by(subject, treatment, measurement) %>%
    mutate(relChange=value/first(value)) %>%
    filter(time != "Pre")
```

Let us repeat our earlier graphs now. We will adjust the y-axis to show relative increase as a percent (more about this stuff in the advanced graphics session):

```
ggplot(relativedata) +
    aes(x=time, y=relChange-1) +
    geom_path(aes(group=subject, color=subject)) +
    facet_grid(measurement~treatment) +
    labs(y="relative change from pre") +
    guides(color=FALSE) +
    scale_y_continuous(labels = scales::percent)

ggplot(relativedata) +
    aes(x=time , y=relChange-1, group=treatment, col=treatment) +
    stat_summary(geom="pointrange", fun.data=mean_se, fatten=3, alpha=0.6,
        position=position_dodge(width=0.1)) +
    geom_line(stat="summary", fun.y=mean, position=position_dodge(width=0.1)) +
    facet_grid(measurement~.) +
    labs(y="relative change from pre") +
    scale_y_continuous(labels = scales::percent)
```

Practice

You likely have a dataset of your own that you might want to work with to practice the above. If you don't, then here's a possible project (make sure to start a new project either way):

1. The GapMinder[6] website has a wealth of information on each country over many years. We will suggest some analyses to perform but you are free to pursue a different analysis. There is actually an R package that can load a lot of this data in more ready form, but we will use this as an opportunity to access data from the web and from multiple sources.

   - Start by creating a new project. You may also want to start an RMarkdown document to keep your work organized.
   - Next we will access the data sets from the website, https :// www.gapminder.org/data/. On that page you can see a list of datasets, and you can use the search box at the upper right of the first table to search for a variable of interest (we will discuss what those are for us in a moment). You have two options: You can either download the file and include it in your project, or you can use the file's URL directly into R. We will recommend the first approach in this instance, but keep in mind that the second path is possible.

2. We will use the `life expectancy ( years )` variable.

   - Search for the corresponding dataset, and download the file. Then upload the file into your project directory (you may want to rename it to something simpler), and import its first sheet into the project by clicking the uploaded file.
   - Start by using `rename` to rename the first variable to `country` or something similar.
   - Use `gather` to turn all the different year variables into a single year. The resulting dataset should have three columns: country, year, and "life expectancy". You can use –`country` as the column specification; this will tell `gather` to use all columns except for `country`. Also make sure to check the documentation and include a setting to "remove the missing values".

   One thing we have not covered before is using variable names with spaces in them, like `life expectancy` in this example. In many places, you will need to enclose them in backticks to make that work, like so: ‘ `life expectancy`‘

   - Notice that the `year` variable in the result is treated as "character". We really would prefer it to be treated as numeric, so add a `convert=TRUE` parameter to the `gather` function, which will tell the system to try to convert the key column (`year`) to whatever is more meaningful given what the values are. Look at the documentation of `gather` for details. (We could also have used a `mutate` step with the `as.numeric` function to post-process the `year` column.)
   - Once you are satisfied that the output looks as expected, store it in some variable (e.g. maybe named `lifeExpectancy` ?).

3. We will now add another variable, the per-capita income. The variable you want is called "Income per person (GDP/capita, PPP$ inflation-adjusted)". Follow much the same steps as in the previous point. You can call the values column whatever you want of course, but we will be using simply `income` for it.

---

[6]https://www.gapminder.org

4. We will also want to get the total population of each country, stored in a dataset named "Population, total". Use similar steps for it.

   (Optional step) If you are feeling adventurous, you can pursue the possibility of writing a function to make these steps easier, since we more or less repeated the same steps three times.

5. Now we want to bring these three datasets together, by matching entries on the country and year. The inner_join function can help you do that; make sure to look at its documentation, or the dplyr cheatsheet. Store the result in a variable (we called it gapdata).

   You'll need to do this in two steps, as inner_join only joins two datasets at a time. So you'll start with one dataset, pipe it into an inner_join with the second dataset, and then pipe the result into another inner_join with the third dataset. So your code might look something like this:

```
gapdata <- lifeExpectancy %>%
    inner_join(income, by = ....) %>%
    inner_join(population, by = ....)
```

   You have a choice to make here, actually, inner_join vs full_join. If you use full_join, you will include entries for country and year combinations where one of the two values of income and life expectancy is missing. This gives us more data to work with, but we might end up having to throw away some values later on anyway when trying to do graphs or analysis and so on. We will use inner_join instead, which will only keep the entries where both the income and the life expectancy are known.

   If you have done this correctly, you should be seeing about 15, 000 observations on $5% variables.

Some questions to try out

This ends the "data tidying part", except for one more challenging addition that we'll discuss after these problems. We can now do some analyses on these numbers.

1. Use arrange (maybe along with head or tail) to reorder the rows in the gapdata and discover what the most recent year in the data is. You can also find that out from the viewer, or you can use summarize or pluck along with the max function.
2. Use filter to consider only the subset of the data that relates to the most recent year, then arrange to order the data according to income, and find the countries with the 5 highest and 5 lowest per capita incomes.
3. For each country, determine how many rows of data we have for that country, and visualize the results with an appropriate graph.
4. Determine for each year the country with the highest per-capita income. You can do this with a group_by and a filter.
5. Compute the total world population per year.
6. Compute total income for each row by multiplying together the population and income variables, and add that as a new variable.

A challenge: Matching Country Names

Warning: This section is somewhat challenging, as it attempts to reconcile differences in country names between different datasets. The end result is simply to have a region variable associated with each country.

Feel free to skip this on a first go.

We will incorporate a variable for the various regions of the world. As this variable does not appear to be available in the GapMinder data, we will need to get it from another source. In this instance, the World Bank seems to include it in some of their available datasets:

http://databank.worldbank.org/data/download/WDI_excel.zip

- Download the file and extra the contained Excel file, then upload it to your project folder.
- The second sheet, named country, contains the information we need. Load that into a variable, we called it wdi.
- Since we only care about the country name and region, use select to only choose the four columns Short Name, Table Name, Long Name, Region.

The challenging part now is to match the country names between the two datasets, which can be different due to different conventions on naming some of the countries. For instance, let us start by making a list of all the countries in gapdata and in the Table Name entry of the wdi data, and look at their difference:

```
gapCountries <- gapdata %>% pluck("country") %>% unique() %>% sort()
wdiCountries <- wdi %>% pluck("Table Name") %>% unique() %>% sort()
setdiff(gapCountries, wdiCountries)    # names in gapCountries but not wdiCountries
```

You should see a number of countries in that list, meaning that they don't show up as such in wdiCountries, but that really must be somewhere there, maybe with a different representation. We could also try to use the Short Name column instead, but if you do that you'll notice that it too is missing some countries.

Before we move on, you might want to View the wdi data and search for those missing countries in the data view, to see how they show up.

We could of course simply edit the Excel sheet and change the country names. But then we may need to do that each time we want to update the data from the originals. Manually changing the Excel sheets is unsustainable in the long run. So we will look for a programmatic way to do it. Here are the steps we will take:

- We first create a new column in the wdi data to hold either the Table Name or the Short Name, whichever one actually matches the country name in the gapData. We will call this new column gapName.
- Next we will find out which countries in gapdata do not have a corresponding entry in the gapName column. Hopefully that will be a small list. We will then manually create a small dataset that contains an appropriate matching name for each of those countries, by inspecting the names that exist in the Table Name column.
- Then we will join this dataset with the wdi to fill in the remaining gapName entries.
- Lastly, we will join our gapdata with this updated wdi dataset in order to get the region names into the gapdata.

So let's start with the first part, adding a new column in wdi. We can use the case_when method for this: It tries a series of alternatives in turn and uses the first one that matches:

```
wdi <- wdi %>% mutate(gapName = case_when(
  'Table Name' %in% gapCountries    ~ 'Table Name',
  'Short Name' %in% gapCountries    ~ 'Short Name',
  TRUE                              ~ as.character(NA)))
setdiff(gapCountries, wdi$gapName)
```

10

So there are 12 countries in the gapdata which do not yet have a matching gapName in the wdi data. Let's make a small dataset that contains these countries and their appropriate wdi entries. We will use the tribble command for that, which allows us to nicely align the values vertically. We used the search feature in the data view in order to find what the appropriate names would be:

```
countryNames <- tribble(
  ~gapName, ~'Short Name',
  "Bahamas", "The Bahamas",
  "Cape Verde", "Cabo Verde",
  "Cote d'Ivoire", "Côte d'Ivoire",
  "Gambia", "The Gambia",
  "Hong Kong, China", "Hong Kong SAR, China",
  "Lao", "Lao PDR",
  "Macao, China", "Macao SAR, China",
  "North Korea", "Dem. People's Rep. Korea",
  "Sao Tome and Principe", "São Tomé and Principe",
  "South Korea", "Korea",
  "Syria", "Syrian Arab Republic"
)
```

We could not find an entry for Taiwan in the wdi data. We will need to assign its region manually after we do all the needed joins.

Now we want to join the countryNames dataset into the wdi data. This is a bit tricky: A join sounds like the right thing to do, with Short Name being the common variable. But the way R does this is to create two separate columns, one for the gapNames from wdi and another for those from countryNames. We need to use coalesce along with mutate to merge the two together (you mean want to check out the coalesce function's documentation and possibly look at the result with just the full_join to see why the mutate step is needed).

```
wdi2 <- wdi %>% full_join(countryNames, by="Short Name") %>%
  mutate(gapName = coalesce(gapName.x, gapName.y))
setdiff(gapCountries, wdi2$gapName)
```

You should now see the setdiff method return only Taiwan. Good, that means each value in gapCountries is now accounted for in wdi. We will still need to deal with Taiwan.

Now, we want to merge in the region values from wdi into our gapdata set. This will simply be a join, matching the country to the gapName. But we only want to bring in the region variable, so before joining we will select only a few columns from wdi2 (left_join here tells it to ignore entries in wdi2 that don't have a matching value in gapdata but to include entries in gapdata that don't have matching value in wdi2. This way Taiwan is not left out):

```
gapdata <- gapdata %>%
    left_join(wdi2 %>% select(gapName, Region),
        by=c("country"="gapName"))
# Now some cleanup and adding Taiwan
gapdata <- gapdata %>%
    rename(region="Region") %>%
    mutate(region=replace(region, country=="Taiwan", "East Asia & Pacific")) %>%
    mutate(region=factor(region))
```

The by part of the left_join method might look different than before. This is because the two columns we wish to join on have different names in the two datasets. The by=c("country"="gapName") syntax tells it to match the country variable in the first dataset to the gapName variable in the second dataset.

Notice the first mutate line above. It tells it to replace the region variable by using whatever was in the region variable before but whenever the country is Taiwan it should set the region to be East Asia and Pacific.

Phew, and with that, we should be done! In order to check, try to use group_by region and summarize using n_distinct to count each country only once. If done correctly, this should give you a reasonable count of countries for each region (for example, 3 countries in North America).

As practice, you can now compute total populations for each region of the world for each year.

Addendum: Saving data

Phew, after all this work we should save this dataset for future use. We can of course save it as a csv file, which would be convenient if you wanted to share it with non-R users. We can do that with the write_csv and write_excel_csv methods. You should prefer these methods to the similarly named write.csv and write.excel, and you can look at their documentation for details. An example call might be:

```
getwd()      # Check to see what the current directory is. This is where files will be saved
write_csv(gapdata, "gapdata.csv", na="")
```

Another option is to save the data set as an R object. You can in fact save any kind of R object, for instance model fits, summary results etc. You can do this with the save command, and the objects can be then loaded back via the load command:

```
save(gapdata,  file="gapdata.RData")
```

It is conventional to use RData for the extension of such a file.

We should point out however, that sharing data this way does not show the work that was used to produce those data, compared to providing an RMarkdown file or an R script, which both show all the processing steps.