Lab 1: Introduction to R and RStudio

Overall Goals

- Start RStudio.
- Learn the different screen areas in RStudio.
- Load a built-in dataset.
- Look at the help for the dataset.
- View the dataset in tabular form.
- Perform basic operations on dataset
  - Number of rows
  - Sort rows based on column
  - Simple graphs
  - Simple numerical summaries

Specific R commands learned

- data for loading built-in data.
- ? for asking for information on a function or data.
- View for a tabular look at dataset.
- nrow for number of rows.
- ...

Lab 1

- Start RStudio (TODO: weblink)
- Panes:
  - Console
  - Environment/History
  - Outputs/Help/File
  - Document (Later)
- Create new project for counties
- Use "library(...)", "data(...)" to load the package that has the counties data (need to create it).
- (Not for student use: read . table (" https :// skiadas . github . io / AppliedStatsCourse / site / datasets /countyComplete.txt ", sep=)
- See the counties set on the right, click on it
- Data view comes up -> document pane (also View command)
- Sort view to find 3 most/least populous counties based on 2010
- How many counties? (3143, on left side)
- number of counties per state:
  - Numerically: tally (~ state , data=counties)
    * Use upper arrow to repeat, then add: %>% sort()
    * Then add: %>% barchart()
  - Graphically: bargraph(~state, data=counties, horizontal = TRUE)

- * zoom graph
- Population of each state:
  - sum(~pop2010|state, data=counties) %>% sort() %>% barchart()
  - Maybe skip for later?
- Histogram of population by county:
  - histogram(~pop2010, data=counties)
  - histogram(~pop2010, data=counties %>% filter(pop2010 <= 2e6))
  - Add "breaks = 40"
  - favstats (~pop2010, data=counties)
- Histogram of percent of female population:
  - histogram(~female, data=counties, breaks=40)
  - favstats (~female, data=counties)
  - bwplot(state ~female, data=counties)
    - * Find lowest median, highest median
    - * Find state with lowest outlier
  - Work with black and/or asian states
- Compare populations 2000, 2010:
  - xyplot(pop2010~pop2000, data=counties)
  - xyplot(pop2010~pop2000, data=counties %>% filter(pop2000 < 2e6))
  - tally (~name, data=counties) %>% sort() %>% tail(10)
- County names:
  - tally (~name, data=counties) %>% sort() %>% tail(10)
- Union counties:
  - counties %>% filter(name=="Union County") %>% select(state)
  - summary

Lab 3 or something: Introduce markdown

TODO: Consider using the %in% operator somewhere

To find percent of whites in state: - white count per county - add white counts across states - divide white count by state pop count

counties %>% transmute(state = state, pop = pop2010, whitepop = pop2010 * white / 100) %>% group_by(state) %>% summarize(pop = sum(pop), whitepop = sum(whitepop)) %>% transmute(state = state, white = whitepop/pop * 100) %>% arrange(white)

Easier: sum(white*pop2010/100~state, data=counties)/sum(pop2010~state, data=counties) 100

sum(~black*pop2010/100, data=counties %>% filter(!is.na(black)))/sum(~pop2010, data=counties %>% filter(!is.na(black)))